DCAT extension benchmarks: evidence and reference layers

Historical snapshot: 7 October 2026 | All six reconstructed corpora are partial

a Two evaluation entry points

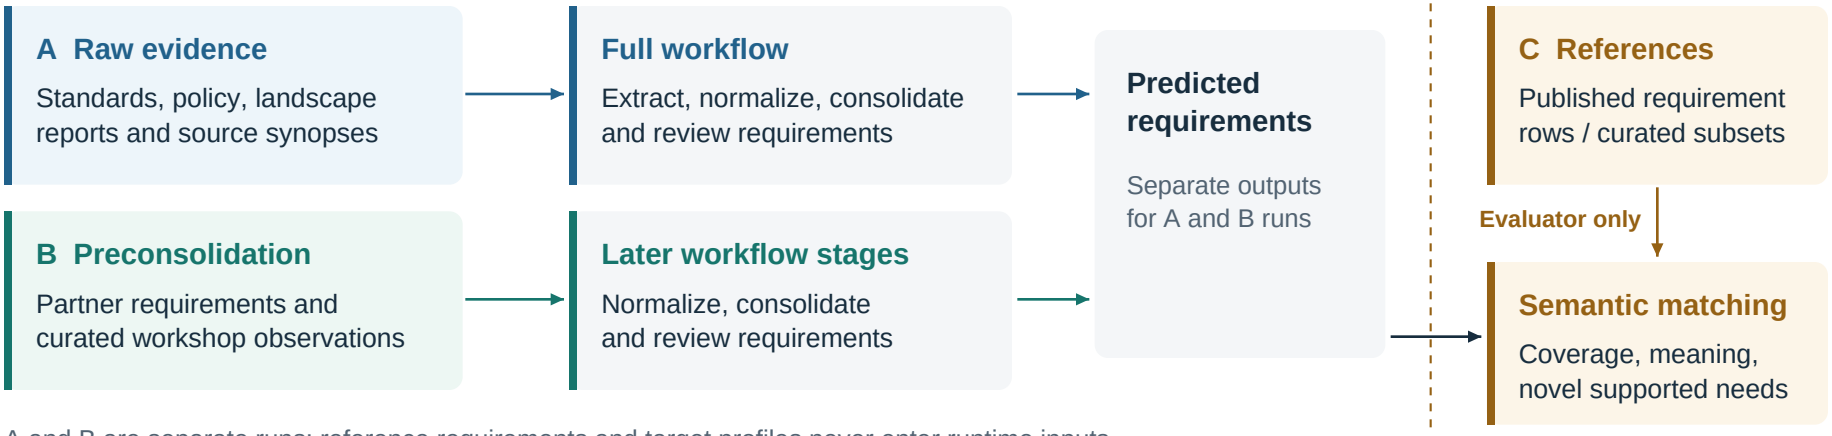

A and B are separate runs; reference requirements and target profiles never enter runtime inputs.

b What is available in this snapshot

Benchmark	A Raw inputs	B Preconsolidation	C Reference scope
mobilityDCAT-AP Primary case	16 files: policy, DCAT/INSPIRE, FAIR, SPRINT and CMC	3 files: 60 partner rows + review / expert observations	40 published Table 1 rows All IDs; curator paraphrases
HealthDCAT-AP Heterogeneous sources	5 files: landscape analysis, EHDS proposal, GDPR, DCAT	1 file: 8 curated TWG observations; partial subset	16 curated rationale entries Original D6.2 also retained
StatDCAT-AP Standards-derived	2 files: DCAT and Data Cube Use-case reconstruction excluded*	Unavailable	7 requirement / resolution pairs Scoped selection from 2016 report
GeoDCAT-AP Standards alignment	5 files: INSPIRE legislation, public guidance and DCAT	Unavailable	7 curated alignment constraints ISO standards are citation-only
JRC research Small sanity case	3 files: data policy, DCAT, DataCite field synopsis	Unavailable	5 enumerated core needs Curator paraphrases
EPOS-DCAT-AP Exploratory only	No scored raw run enabled Public context is available	Unavailable	No reconstructed requirement set Frozen specification + context

Counts refer to allowlisted files in A/B and reference entries in C; they are not directly comparable.
*StatDCAT use cases are available as a separate retrospective variant. Curated entries await independent review.
Source: frozen corpus manifests, runtime allowlists and reference inventories. No workflow results are shown.